

# Linux 0.11进程调度

张延慈

# 输出进程的信息

- schedule is called by do\_exit.
- Maybe 15 is the longest counter and 3 is the chosen one?
- The chosen task:
- 3: pid=3, state=0, counter = 15, father=1
- jiffies is 4
- all tasks as follows:
- 0: pid=0, state=1, counter = 14, father=-1
- jiffies is 4
- 1: pid=1, state=0, counter = 14, father=0
- jiffies is 4
- 2: pid=2, state=3, counter = 13, father=1
- jiffies is 4
- 3: pid=3, state=0, counter = 15, father=1
- jiffies is 4

```
struct task_struct {  
  
    long state; //表示进程的状态, -1表示不可执行, 0表示可执行, >0表示停止  
    long counter; /* 运行时间片,以jiffies递减计数 */  
    long priority; /* 运行优先数,开始时, counter = priority, 值越大,表示优先数越高,等待时  
    long signal; /* 信号.是一组位图,每一个bit代表一种信号. */  
    struct sigaction sigaction[32]; /* 信号响应的数据结构, 对应信号要执行的操作和标志信.  
    long blocked; /* 进程信号屏蔽码(对应信号位图) */  
    /* various fields */  
    int exit_code; /* 任务执行停止的退出码,其父进程会取 */  
    unsigned long start_code, end_code, end_data, brk, start_stack; /* start_code代码段地  
    end_data代码长度+数据长度(byte), brk总长度(byte), start_stack堆栈段地址 */  
    long pid, father, pgrp, session, leader; /* 进程号,父进程号,父进程组号,会话号,会话头(发  
    unsigned short uid, euid, suid; /* 用户id号,有效用户id号,保存用户id号*/  
    unsigned short gid, egid, sgid; /* 组标记号(组id),有效组id,保存的组id */  
    long alarm; /* 报警定时值(jiffies数) */  
    long utime, stime, cutime, cstime, start_time; /* 用户态运行时间(jiffies数),  
    系统态运行时间(jiffies数),子进程用户态运行时间,子进程系统态运行时间,进程开始运行时刻 */  
    unsigned short used_math; /* 是否使用了协处理器 */  
    /* file system info */  
    int tty; /* 进程使用tty的子设备号. -1表示没有使用 */  
    unsigned short umask; /* 文件创建属性屏蔽位 */  
    struct m_inode * pwd; /* 当前工作目录i节点结构 */  
    struct m_inode * root; /* 根目录i节点结构 */  
    struct m_inode * executable; /* 执行文件i节点结构 */  
    unsigned long close_on_exec; /* 执行时关闭文件句柄位图标志. */  
    struct file * filp[NR_OPEN];  
    /* 文件结构指针表,最多32项. 表项号即是文件描述符的值 */  
    struct desc_struct ldt[3];  
    /* 任务局部描述符表. 0-空, 1-cs段, 2-Ds和Ss段 */  
    struct tss_struct tss; /* 进程的任务状态段信息结构 */
```

# 输出调用进程调度函数的位置

- 在exit.c和schedule.c中， schedule.c被调用

```
int sys_pause(void)
{
    current->state = TASK_INTERRUPTIBLE;
    log("Schedule is called by sys_pause.\n");
    schedule();
    log("\n");
    return 0;
}

void sleep_on(struct task_struct **p)
{
    struct task_struct *tmp;

    if (!p)
        return;
    if (current == &(init_task.task))
        panic("task[0] trying to sleep");
    tmp = *p;
    *p = current;
    current->state = TASK_UNINTERRUPTIBLE;
    log("Schedule is called by sleep_on.\n");
    schedule();
    log("\n");
    *p = tmp;
    if (tmp)
        tmp->state=TASK_RUNNING;
}
```

```

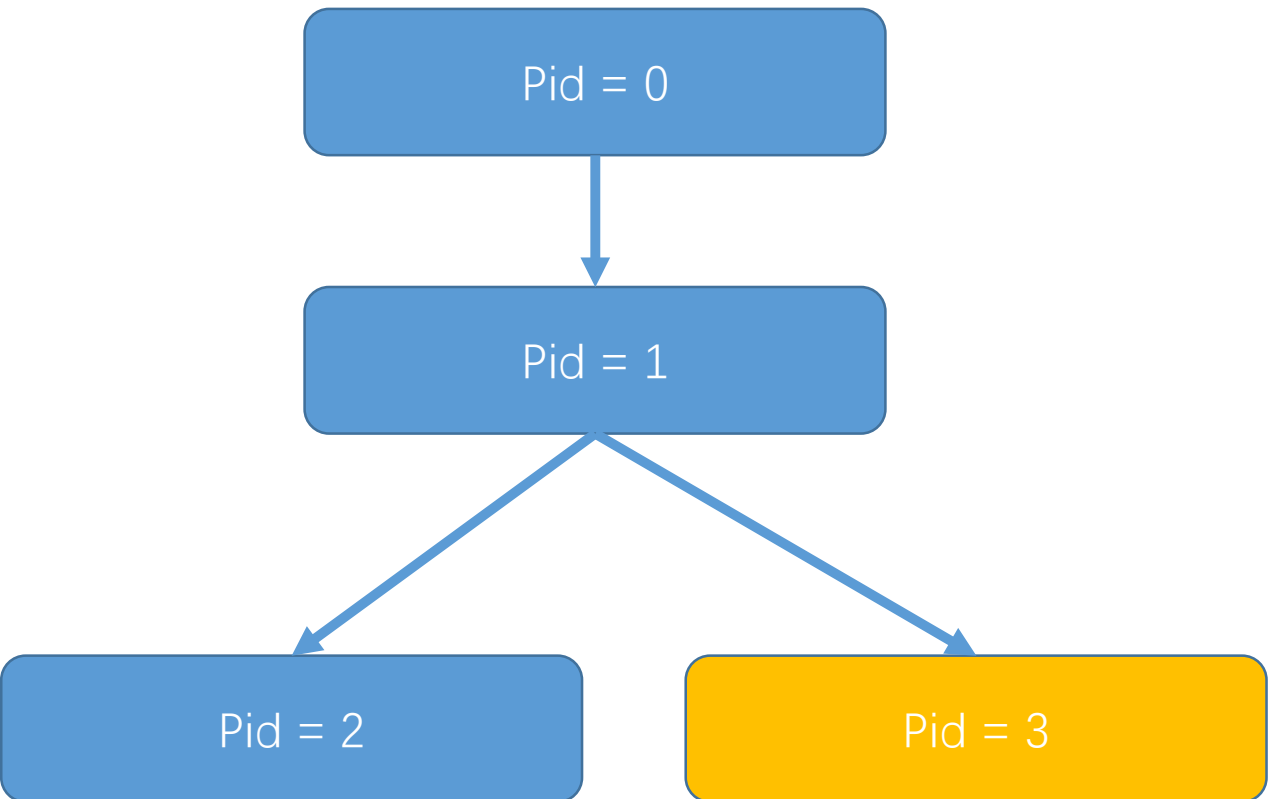
void show_task(int nr, struct task_struct * p)
{
    //int i,j = 4096-sizeof(struct task_struct);
    //printk("%d: pid=%d, state=%d, ",nr,p->pid,p->state);
    log("%d: pid=%d, state=%d, counter = %d, father=%d \njiffies is %d\n\n",nr,p->pid,p->state, p->counter, p->father, jiffies);
    //i=0;
    //while (i<j && !((char *)(p+1))[i])
    //    i++;
    //printk("%d (of %d) chars free in kernel stack\n\r",i,j);
    //log("%d (of %d) chars free in kernel stack\n",i,j);
}

void show_stat(void)
{
    int i;

    for (i=0;i<NR_TASKS;i++)
        if (task[i])
            show_task(i,task[i]);

    log("\n");
}

```



- 3: pid=3, state=0, counter = 15, father=1
- jiffies is 4
  
- all tasks as follows:
- 0: pid=0, state=1, counter = 14, father=-1
- jiffies is 4
- 1: pid=1, state=0, counter = 14, father=0
- jiffies is 4
- 2: pid=2, state=3, counter = 13, father=1
- jiffies is 4
- 3: pid=3, state=0, counter = 15, father=1
- jiffies is 4

---

Schedule is called by sys\_pause.

Maybe 15 is the longest counter and 1 is the chosen one?

The chosen task:

1: pid=1, state=0, counter.=15, father=0.

jiffies is 1

all tasks as follows:

0: pid=0, state=1, counter.=14, father=-1.

jiffies is 1

1: pid=1, state=0, counter.=15, father=0.

jiffies is 1

Maybe 15 is the longest counter and 2 is the chosen one?

The chosen task:

2: pid=2, state=0, counter.=15, father=1.  
jiffies is 1

all tasks as follows:

0: pid=0, state=1, counter.=14, father=-1.  
jiffies is 1

1: pid=1, state=1, counter.=15, father=0.  
jiffies is 1

2: pid=2, state=0, counter.=15, father=1.  
jiffies is 1

schedule is called by do\_exit.

Maybe 15 is the longest counter and 3 is the chosen one?

The chosen task:

3: pid=3, state=0, counter.=.15, father=1.

jiffies is 4

all tasks as follows:

0: pid=0, state=1, counter.=.14, father=-1.

jiffies is 4

1: pid=1, state=0, counter.=.14, father=0.

jiffies is 4

2: pid=2, state=3, counter.=.13, father=1.

jiffies is 4

3: pid=3, state=0, counter.=.15, father=1.

jiffies is 4



Schedule is called by sys\_pause.

The chosen task:

0: pid=0, state=1, counter = 0, father=-1.

jiffies is 25